

# SQL Overview

- Query capabilities
  - SELECT-FROM-WHERE blocks,
  - Basic features, ordering, duplicates
  - Set ops (union, intersect, except)
  - Aggregation & Grouping
  - Nested queries (correlation)
  - Null values

# Nested queries

- A **nested query** is a query with another query embedded within it.
- The embedded query is called the **subquery**.
- The subquery usually appears in the WHERE clause:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                   FROM Reserves R
                   WHERE R.bid = 103 )
```

(Subqueries also possible in FROM or HAVING clause.)

# Conceptual evaluation, extended

- For each row in cross product of outer query, evaluate the WHERE clause conditions, (re)computing the subquery.

```
SELECT  S.sname  
FROM    Sailors S  
WHERE   S.sid IN ( SELECT R.sid  
                   FROM Reserves R  
                   WHERE R.bid = 103 )
```


equivalent to:

```
SELECT  S.sname  
FROM    Sailors S, Reserves R  
WHERE   S.sid=R.sid AND R.bid=103
```

# Correlated subquery

- If the inner subquery depends on tables mentioned in the outer query then it is a **correlated subquery**.
- In terms of conceptual evaluation, we must recompute subquery for each row of outer query.

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS ( SELECT *
                  FROM    Reserves R
                  WHERE   R.bid = 103
                  AND R.sid = S.sid )
```



Correlation

# Set-comparison operators

- Optional NOT may precede these:
  - **EXISTS R** -- true if R is non-empty
  - **attr IN R** -- true if R contains attr
  - **UNIQUE R** -- true if no duplicates in R
- For arithmetic operator **op** {<, <=, =, < >, >=, >}
  - **attr op ALL R** -- all elements of R satisfy condition
  - **attr op ANY R** -- some element of R satisfies condition

**IN** equivalent to **= ANY**

**NOT IN** equivalent to **< > ALL**

# Example

- Find the sailors with the highest rating

```
SELECT S.sid, S.name  
FROM   Sailors S  
WHERE  S.rating >= ALL (SELECT S2.rating  
                        FROM Sailors S2 )
```

# i-clicker

What is the result of these two queries on the Sailors relation?

**SAILORS**

sid	sname	rating	age
29	brutus	1	33
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5
22	dustin	7	45
64	horatio	7	35
31	lubber	8	55.5
32	andy	8	25.5
74	horatio	9	35
58	rusty	10	35
71	zorba	10	16

**A**

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ANY (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Horatio')
```

**B**

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ALL (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Horatio')
```

# i-clicker

What is the result of these two queries on the Sailors relation?

A1	A2
sid	sid
31	31
32	32
74	74
58	58
71	71

B1	B2
sid	sid
22	74
64	58
31	71
32	
74	
58	
71	

C1	C2
sid	sid
31	58
32	71
74	
58	
71	

D1	D2
sid	sid
74	74
58	58
71	71

**E** Honestly, I have no idea.



Answer on next slide

# i-clicker

What is the result of these two queries on the Sailors relation?

- Find sailors whose rating is higher than **some** sailor named Horatio.

A

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ANY (SELECT S2.rating
                       FROM Sailors S2
                       WHERE S2.name = 'Horatio')
```

- B
- Find sailors whose rating is higher than **all** sailors named Horatio.

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ALL (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Horatio')
```

# i-clicker

What is the result of these two queries on the Sailors relation?

A1	A2
sid	sid
31	31
32	32
74	74
58	58
71	71

B1	B2
sid	sid
22	74
64	58
31	71
32	
74	
58	
71	

C1	C2
sid	sid
31	58
32	71
74	
58	
71	

D1	D2
sid	sid
74	74
58	58
71	71

Find boats **not** reserved by sailor with  
sid = 100.

- B: all boats
- R: boats reserved by sailor with sid=100
- B – R is what we want.

```
SELECT B.bid  
FROM Boats B  
WHERE B.bid NOT IN (SELECT R.bid  
                     FROM Reserves R  
                     WHERE R.sid = 100 );
```

# Existential conditions

- Find the names of sailors who have reserved **some** boat
- (i.e. *there exists* a boat they reserved)

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid
```

- Existential conditions are natural and easy in SQL.

# Universal conditions

- Find the names of sailors who have reserved **all** boats.
- (i.e. *for each* boat, they have reserved it.)
- Universal conditions are harder...


```
SELECT S.sname
FROM   Sailors S
WHERE  NOT EXISTS (
           Set of boats not
           reserved by S.sid
        )
```

# Universal conditions

- Find the names of sailors who have reserved **all** boats.

```
SELECT S.sname  
FROM   Sailors S  
WHERE  NOT EXISTS (
```

boats the sailor hasn't reserved  
(we just wrote this query)



```
SELECT B.bid  
FROM Boats B  
WHERE B.bid NOT IN (SELECT R.bid  
                   FROM Reserves R  
                   WHERE R.sid = S.sid )
```

```
)
```

For each sailor, check that there is no boat s/he hasn't reserved.

*Find the names of sailors who reserved a **red** and a **green** boat.*

using INTERSECT

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

**INTERSECT**

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'green'
```

without INTERSECT

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
      AND S.sid IN
      (SELECT S2.sid
       FROM Sailors S2, Reserves R2, Boats B2
       WHERE S2.sid = R2.sid AND R2.bid = B2.bid AND B2.color = 'green' )
```

“Find all sailors who have reserved a red boat and, further, have **sids** that are included in the set of **sids** of sailors who have reserved a green boat.”



# Simulating EXCEPT (set difference)

- Suppose we have tables R(a,b) and S(a,b)
  - What does this query compute?

```
SELECT DISTINCT *  
FROM R  
WHERE (R.a, R.b) NOT IN (SELECT *  
                        FROM S );
```

Can this be expressed without a nested query?

# Query monotonicity

- A query  $Q$  is monotonic if, adding tuples to an instance  $I$  can only grow the result of  $Q(I)$  but never shrink it.
  - Let  $I$  be any database instance.
  - Consider the instance that adds tuple  $t$ :  $I \cup \{t\}$
  - Compare:  $Q(I)$  with  $Q(I \cup \{t\})$
  - If  $Q$  is monotonic, then  $Q(I) \subseteq Q(I \cup \{t\})$

# Monotonicity

- QUESTION: Is this query monotonic?

```
SELECT S.sid, S.name  
FROM   Sailors S  
WHERE  S.rating >= ALL (SELECT S2.rating  
                        FROM Sailors S2 )
```

- FACT: All simple Select-From-Where queries (without subqueries) are monotonic.
- QUESTION: Can the query above be expressed without a nested query?

# SQL Overview

- Query capabilities
  - SELECT-FROM-WHERE blocks,
  - Basic features, ordering, duplicates
  - Set operations (union, intersect, except)
  - Aggregation & Grouping
    - Nested queries (correlation)
  - Null values
  - Outer joins

# NULLS in SQL

- Whenever we don't have a value for a cell, we can put a special value "**NULL**"
- Null can mean many things:
  - *Value does not exists*
  - *Value exists but is unknown*
  - *Value not applicable*
- The schema specifies for each attribute whether it is allowed to be null (*nullable* attribute)

```
CREATE TABLE Sailor (  
    sid NOT NULL  
    ... )
```

# Outer Joins

- In a typical join, tuples of one relation that don't match any tuple from the other relation are omitted from result.
- In an outer join, **tuples without a match can be preserved** in the output.
- Missing values are filled with NULL.

# Outer Joins

- LEFT OUTER JOIN: rows of **left** relation without matching row in right relation appear in result.
- RIGHT OUTER JOIN: rows of **right** relation without matching row in left relation appear in result.
- FULL OUTER JOIN: rows of both relations

Standard joins we have seen so far are called **inner joins** to distinguish them from outer joins.

# Outer Joins

Sailors

sid	sname	rating	age
22	dustin	7	45
31	lubber	8	55.5
58	rusty	10	35

Reserves

sid	bid	day
22	101	10/10
58	103	10/10
58	105	12/1

Sailors LEFT OUTER JOIN Reserves:

sid	sname	rating	age	bid	day
22	dustin	7	45	101	10/10
31	lubber	8	55.5	NULL	NULL
58	rusty	10	35	103	10/10
58	rusty	10	35	105	12/1



# Full Outer Join

(Unless otherwise specified, we assume the **natural** join)

R	A	B	S	B	C
	22	foo		bar	X
	31	bar		wee	Y
	58	wee		yah	Z

Full Outer Join of R and S:

A	B	B	C
22	foo	NULL	NULL
31	bar	bar	X
58	wee	wee	Y
NULL	NULL	yah	Z

# i-clicker

Compute quantities for each of the following joins on R and S:

- (i) Number of tuples in the crossproduct
- (ii) Number of tuples in the natural join
- (iii) Number of tuples in the left outer join
- (iv) Number of tuples in the right outer join
- (v) Number of tuples in the full outer join

**R**

A	B
alfalfa	3
barley	5
chard	2
dandelion	1

**S**

B	C
5	X
3	Y
4	Z

What is the sum: (i) + (ii) + (iii) + (iv) + (v)

- A. 20
- B. 24
- C. 26
- D. 27
- E. 28

Answer on next slide

# i-clicker

Compute quantities for each of the following joins on R and S:

- (i) Number of tuples in the crossproduct = 12
- (ii) Number of tuples in the natural join = 2
- (iii) Number of tuples in the left outer join = 4
- (iv) Number of tuples in the right outer join = 3
- (v) Number of tuples in the full outer join = 5

**R**

A	B
alfalfa	3
barley	5
chard	2
dandelion	1

**S**

B	C
5	X
3	Y
4	Z

What is the sum: (i) + (ii) + (iii) + (iv) + (v)

- A. 20
- B. 24
- C. 26
- D. 27
- E. 28

# Alternative Join Specification in SQL

FROM Sailors S NATURAL JOIN Reserves R
FROM Sailors S JOIN Reserves R ON (S.sid = R.sid)
FROM Sailors S LEFT OUTER JOIN Reserves R ON (S.sid = R.sid)
FROM (Sailors S NATURAL JOIN Reserves R) NATURAL JOIN Boats B

# Database Views

# Views

- A **view** is a relation defined by a query.
- The query defining the view is called the **view definition**

# Two Kinds of Views

Virtual View	Materialized View
the view relation is defined, but not computed or stored.	the view relation is computed and stored in system.
Computed only on-demand (slower at runtime).	Pre-computed offline (faster at runtime).
Always up to date	May have stale data
CREATE VIEW Developers AS SELECT .... FROM ... WHERE	CREATE TABLE Developers AS SELECT .... FROM ... WHERE



# Virtual view example

Person(name, city)

Purchase(buyer, seller, product, store)

Product(name, maker, category)

```
CREATE VIEW Seattle-view AS
```

```
    SELECT buyer, seller, product, store
```

```
    FROM   Person, Purchase
```

```
    WHERE  Person.city = "Seattle" AND  
           Person.name = Purchase.buyer
```

We have a new **virtual** table:

**Seattle-view(buyer, seller, product, store)**

# View Example

We can use the view in a query as we would any other relation:

```
SELECT name, store
FROM Seattle-view, Product
WHERE Seattle-view.product = Product.name AND
      Product.category = "shoes"
```

# Querying a virtual view

```
SELECT name, Seattle-view.store
FROM   Seattle-view, Product
WHERE  Seattle-view.product = Product.name AND
       Product.category = "shoes"
```



“View expansion”

```
SELECT name, Purchase.store
FROM   Person, Purchase, Product
WHERE  Person.city = "Seattle" AND
       Person.name = Purchase.buyer AND
       Purchase.product = Product.name AND
       Product.category = "shoes"
```

# i-clicker

Given this view definition:

```
CREATE VIEW recent_reservations AS  
SELECT S.sid, S.sname, R.sid, R.bid, R.day  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND R.day > 5
```

Compute the result of this query:

```
SELECT *  
FROM recent_reservations V, Boats B  
WHERE B.bid = V.bid AND B.color = red
```

**Sailors**

sid	sname	rating	age
29	brutus	1	33
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5
22	dustin	7	45
64	horatio	7	35
31	lubber	8	55.5
32	andy	8	25.5
74	horatio	9	35
58	rusty	10	35
71	zorba	10	16

**Reserves**

sid	bid	day
22	101	1
22	102	9
22	103	2
22	104	10
31	102	4
31	103	6
31	104	3
64	101	7
64	102	6
74	103	7

**Boats**

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

# i-clicker

Given this view definition:

```
CREATE VIEW recent_reservations AS  
SELECT S.sid, S.sname, R.sid, R.bid, R.day  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND R.day > 5
```

Compute the result of this query:

```
SELECT *  
FROM recent_reservations V, Boats B  
WHERE B.bid = V.bid AND B.color = red
```

How many tuples are there in the result?

- A. 0
- B. 1
- C. 2
- D. 3
- E. 5

Answer on next slide

# i-clicker

Given this view definition:

```
CREATE VIEW recent_reservations AS  
SELECT S.sid, S.sname, R.sid, R.bid, R.day  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND R.day > 5
```

Compute the result of this query:

```
SELECT *  
FROM recent_reservations V, Boats B  
WHERE B.bid = V.bid AND B.color = red
```

How many tuples are there in the result?

- A. 0
- B. 1
- C. 2
- D. **3**
- E. 5

# The great utility of views

- Data independence (virtual)
- Efficient query processing (materialized)
  - materializing certain results can improve query execution
- Controlling access (virtual)
  - Grant access to views only to filter data
- Data integration (virtual)
  - Combine data sources using views



# View-related issues

## 1. View selection

- which views to materialize, given workload

## 2. View maintenance

- when base relations change, (materialized) views need to be refreshed.

## 3. Updating virtual views

- can users update relations that don't exist?

## 4. Answering queries using views

- when only views are available, what queries over base relations are answerable?

# Updating Virtual Views

How can I insert a tuple into a table that doesn't exist?

`Employee(ssn, name, department, project, salary)`

```
CREATE VIEW Developers AS
SELECT name, project
FROM Employee
WHERE department = "Development"
```

If we make the  
following insertion:

```
INSERT INTO Developers VALUES("Joe", "Optimizer")
```

It becomes:

```
INSERT INTO Employee(ssn, name, department, project, salary)
VALUES(NULL, "Joe", "Development", "Optimizer", NULL)
```

# Non-Updatable Views

Person(name, city)

Purchase(buyer, seller, product, store)

```
CREATE VIEW City-Store AS
```

```
    SELECT Person.city, Purchase.store  
    FROM   Person, Purchase  
    WHERE  Person.name = Purchase.buyer
```

How can we add the following tuple to the view?

(“Seattle”, “Nine West”)

We don't know the name of the person who made the purchase;  
cannot set to NULL.

# Another troublesome example

```
CREATE VIEW OldEmployees AS  
  SELECT name, age  
  FROM Employee  
  WHERE age > 30
```

```
INSERT INTO OldEmployees VALUES("Joe", 28)
```

If this tuple is inserted into the view, it won't appear. Allowed by default in SQL!

# Ambiguous updates

Name	group
Alice	fac
Bob	fac
Bob	cv

group	file
fac	foo.txt
fac	bar.txt
cv	foo.txt

Join

view

Alice	foo.txt
Alice	bar.txt
Bob	foo.txt
Bob	bar.txt

Delete ("Alice", "foo.txt")

# Updating views in practice

- Updates on views highly constrained:
  - SQL-92: updates only allowed on single-table views with projection, selection, no aggregates.
  - SQL-99: takes into account primary keys; updates on multiple table views may be allowed.
  - SQL-99: distinguishes between updatable and insertable views