

Evaluation of Relational Operations

CMPSCI 445

Spring 2018

Relational Operations

- ❖ We will consider how to implement:
 - Selection (σ) Selects a subset of rows from relation.
 - Projection (π) Deletes unwanted columns from relation.
 - Join (\bowtie) Allows us to combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
 - Aggregation (SUM, MIN, etc.) and GROUP BY
 - Order By Returns tuples in specified order.
- ❖ After we cover the operations, we will discuss how to *optimize* queries formed by composing them.

Outline


- ❖ **Sorting**
- ❖ Evaluation of joins
- ❖ Evaluation of other operations

Why Sort?

- ❖ A classic problem in computer science!
- ❖ Important utility in DBMS:
 - Data requested in sorted order (e.g., ORDER BY)
 - e.g., find students in increasing *gpa* order
 - Sorting useful for eliminating *duplicates* (e.g., SELECT DISTINCT)
 - *Sort-merge* join algorithm involves sorting.
 - Sorting is first step in *bulk loading* B+ tree index.
- ❖ Problem: sort 100Gb of data with 1Gb of RAM.

2-Way Sort: Requires 3 Buffers

- ❖ Pass 0: Read a page, sort it, write it.
 - only one buffer page is used
- ❖ Pass 1, 2, ..., etc.:
 - three buffer pages used.



Two-Way External Merge Sort

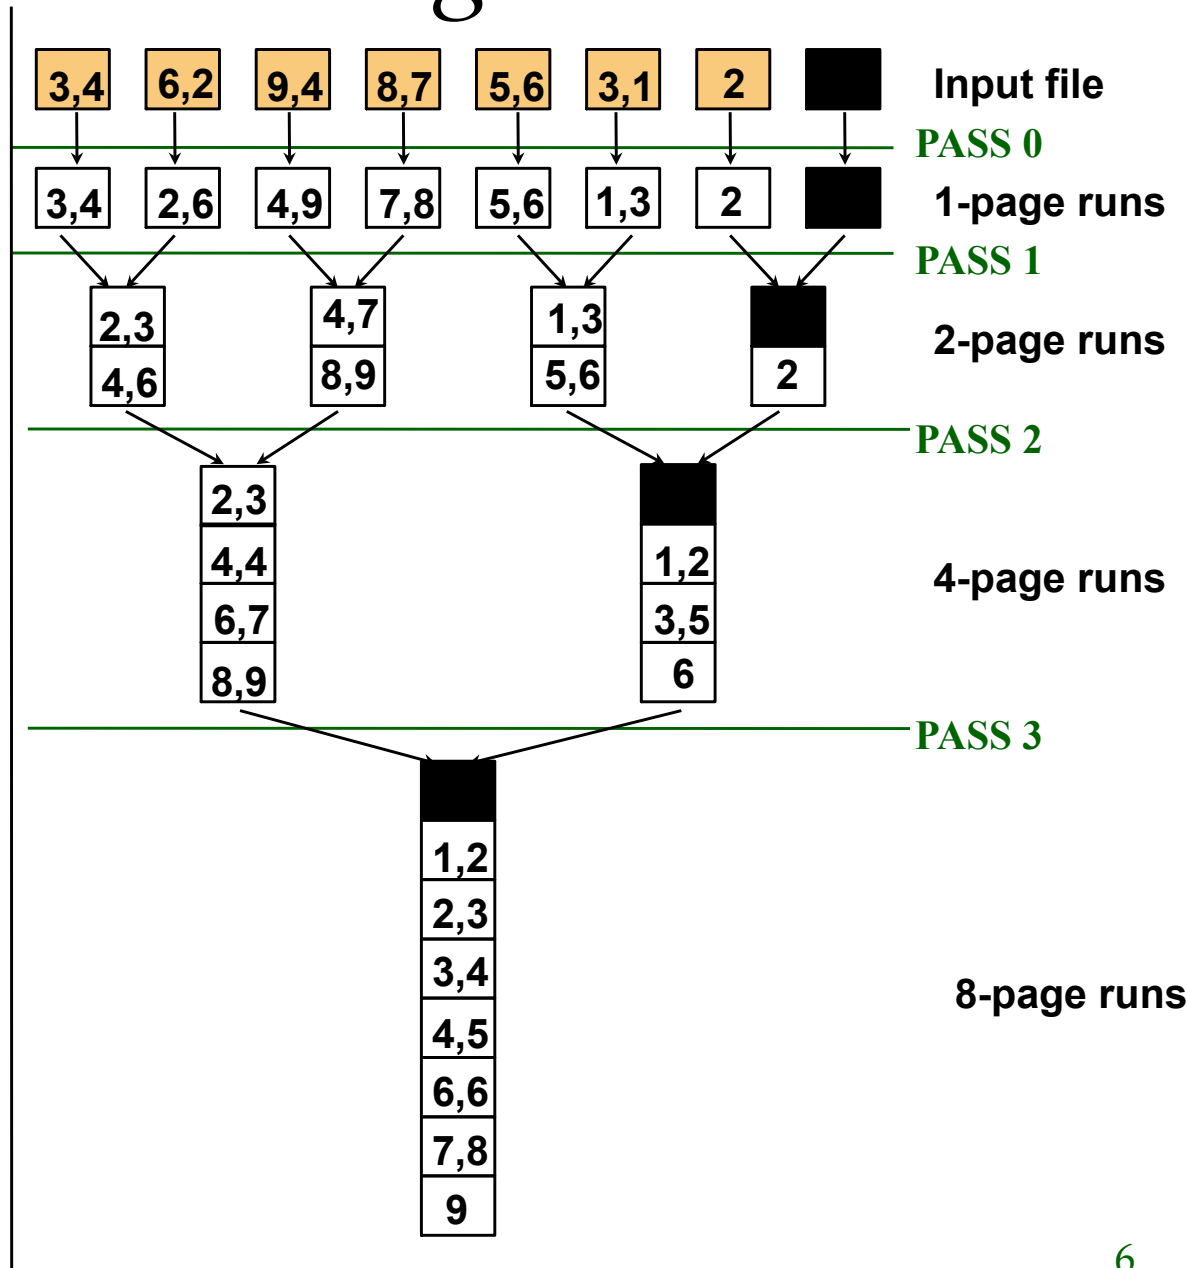
❖ Each pass we read + write each page in file: $2N$.

❖ N pages in the file \Rightarrow the number of passes
 $= \lceil \log_2 N \rceil + 1$

❖ So total cost is:

$$2N(\lceil \log_2 N \rceil + 1)$$


❖ Idea: *Divide and conquer*: sort subfiles and merge



General External Merge Sort

➡ *More than 3 buffer pages. How can we utilize them?*

- ❖ To sort a file with N pages using B buffer pages:
 - Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.
 - Pass 2, ..., etc.: merge $B-1$ runs.



iClicker 1

- ❖ For external merge sort using B buffers, the $(B-1)$ -way merges produce runs of length $(B-1)$.
 - A. True
 - B. False

Answer on next slide

iClicker 1

❖ For external merge sort using B buffers, the $(B-1)$ -way merges produce runs of length $(B-1)$.

A. True

B. False

Cost of External Merge Sort

- ❖ Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- ❖ Cost = $2N * (\text{\# of passes})$
- ❖ E.g., with 5 buffer pages, to sort 108 page file:
 - Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - Pass 2: 2 sorted runs, 80 pages and 28 pages
 - Pass 3: Sorted file of 108 pages

Number of Passes of External Sort

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Don't forget: the full cost of external sort is $2N \times$ (number of passes)

Sorting Records!


- ❖ Sorting has become highly competitive!
 - Parallel sorting is the name of the game ...
- ❖ Datamation sort benchmark: Sort 1M records of size 100 bytes
 - in 1985: 15 minutes
- ❖ New benchmarks proposed:
 - Minute Sort: How many can you sort in 1 minute?
 - Dollar Sort: How many can you sort for \$1.00?

Using B+ Trees for Sorting

- ❖ Scenario: Table to be sorted has B+ tree index on sorting column(s).
- ❖ **Idea:** Can retrieve records in order by traversing leaf pages.
- ❖ *Is this a good idea?*
- ❖ Cases to consider:
 - B+ tree is clustered *Good idea!*
 - B+ tree is not clustered *Could be a very bad idea!*

Clustered B+ Tree Used for Sorting

- ❖ Cost: root to the left-most leaf, then retrieve all leaf pages (if data entries are records)
- ❖ Otherwise, additional cost of retrieving data records: each page fetched just once.




➡ *Always better than external sorting!*

Unclustered B+ Tree Used for Sorting

- ❖ Each data entry contains *rid* of a data record. In general, **one I/O per data record!**

Worse case I/O: pN
 p : # records per page
 N : # pages in file



Summary

- ❖ External sorting is important; DBMS may dedicate part of buffer pool for sorting!
- ❖ External merge sort minimizes disk I/O cost:
 - Pass 0: Produces sorted *runs* of size B (# buffer pages). Later passes: *merge* runs.
 - # of runs merged at a time depends on B .
 - In practice, # of runs rarely more than 2 or 3.
- ❖ Clustered B+ tree is good for sorting; unclustered tree is usually very bad.

iClicker 2

- ❖ Consider scanning the N pages of a relation, searching for records satisfying some property (e.g. age = 21).
 - With 1 input buffer page and 1 output buffer page, this will require N IOs.
- ❖ With B buffer pages in total, how many IO's will it require?
 - A. N
 - B. N/B
 - C. $N/(B-1)$
 - D. $\log_B N$
 - E. $\log_{B-1} N$

Answer on next slide

iClicker 2

- ❖ Consider scanning the N pages of a relation, searching for records satisfying some property (e.g. age = 21).
 - With 1 input buffer page and 1 output buffer page, this will require N IOs.
- ❖ With B buffer pages in total, how many IO's will it require?
 - A. N
 - B. N/B
 - C. $N/(B-1)$
 - D. $\log_B N$
 - E. $\log_{B-1} N$

iClicker

- ❖ What's going to happen Wednesday
 - A. Delayed opening (10 am)
 - B. Campus will close early (12pm)
 - C. Full all-day closure

iClicker

- ❖ What's going to happen Thursday
 - A. Delayed opening (10 am)
 - B. Other delayed opening (12pm?)
 - C. Campus will close early (12pm)
 - D. Full all-day closure

Outline

- ❖ Sorting
- ❖ Evaluation of joins
- ❖ Evaluation of other operations

EXPLAIN in PgAdmin3

Query - simple_facebook on master@cs445hw.cxbepp9iqfon.us-east-1.rds.a

SQL Editor Graphical Query Builder **Click on this button**

Previous queries

```
select * from friend, likes where friend.fid=likes.fid;
```

Output pane

Data Output **Explain** Messages History

The diagram illustrates the execution plan for the query. It shows two input tables, 'friend' and 'likes', each represented by a grid icon. Arrows from both tables point to a 'Hash' operator, which is also represented by a grid icon. From the 'Hash' operator, an arrow points to a 'Hash Join' operator, represented by a grid icon with a green arrow. The 'Hash Join' operator then points to the final output, represented by a grid icon.

OK. Unix Ln 1

----- TPCB Q1: Single Relation -----

```
select
    l_returnflag,
    l_linestatus,
    sum(l_quantity) as sum_qty,
    sum(l_extendedprice) as sum_base_price,
    sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
    sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
    avg(l_quantity) as avg_qty,
    avg(l_extendedprice) as avg_price,
    avg(l_discount) as avg_disc,
    count(*) as count_order
from
    lineitem
where
    l_shipdate <= date '1998-12-01' - interval '82' day
group by
    l_returnflag,
    l_linestatus
order by
    l_returnflag,
    l_linestatus;
```

QUERY PLAN

Sort (cost=4306256.71..4306256.73 rows=6 width=36)

Sort Key: l_returnflag, l_linestatus

-> HashAggregate (cost=4306256.53..4306256.63 rows=6 width=36)

Group Key: l_returnflag, l_linestatus

-> Seq Scan on lineitem (cost=0.00..1936078.65 rows=59254447 width=36)

Filter: (l_shipdate <= '1998-09-10 00:00:00'::timestamp without time zone)

(6 rows)

----- TPCH Q3: 3-way join-----

```
select
    l_orderkey,
    sum(l_extendedprice * (1 - l_discount)) as revenue,
    o_orderdate,
    o_shippriority
from
    customer,
    orders,
    lineitem
where
    c_mktsegment = 'BUILDING'
    and c_custkey = o_custkey
    and l_orderkey = o_orderkey
    and o_orderdate < date '1995-03-22'
    and l_shipdate > date '1995-03-22'
group by
    l_orderkey,
    o_orderdate,
    o_shippriority
order by
    revenue desc,
    o_orderdate;
```

QUERY PLAN

Sort (cost=4253633.68..4261644.36 rows=3204270 width=28)

Sort Key: (sum((lineitem.l_extendedprice * (1::double precision - lineitem.l_discount)))),
orders.o_orderdate

-> GroupAggregate (cost=3665934.82..3754052.25 rows=3204270 width=28)

Group Key: lineitem.l_orderkey, orders.o_orderdate, orders.o_shippriority

-> Sort (cost=3665934.82..3673945.50 rows=3204270 width=28)

Sort Key: lineitem.l_orderkey, orders.o_orderdate, orders.o_shippriority

-> Hash Join (cost=693200.74..3166353.39 rows=3204270 width=28)

Hash Cond: (lineitem.l_orderkey = orders.o_orderkey)

-> Seq Scan on lineitem (cost=0.00..1936078.65 rows=32176879 width=20)

Filter: (l_shipdate > '1995-03-22'::date)

-> Hash (cost=667234.93..667234.93 rows=1493745 width=12)

-> Hash Join (cost=60175.62..667234.93 rows=1493745 width=12)

Hash Cond: (orders.o_custkey= customer.c_custkey)

-> Seq Scan on orders (cost=0.00..455546.00 rows=7311526 width=16)

Filter: (o_orderdate < '1995-03-22'::date)

-> Hash (cost=55147.00..55147.00 rows=306450 width=4)

-> Seq Scan on customer (cost=0.00..55147.00 rows=306450 width=4)

Filter: (c_mktsegment = 'BUILDING'::bpchar)

(18 rows)

Some Common Techniques

- ❖ Algorithms for evaluating relational operators use some simple ideas extensively:
 - **Indexing:** Can use WHERE conditions to retrieve small set of tuples (selections, joins)
 - **Iteration:** Sometimes, faster to scan all tuples even if there is an index. (And sometimes, we can scan the data entries in an index instead of the table itself.)
 - **Partitioning:** By using sorting or hashing, we can partition the input tuples and replace an expensive operation by similar operations on smaller inputs.

** Watch for these techniques as we discuss query evaluation!*

Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: date, rname: string)

❖ Reserves:

- Each tuple is 40 bytes long,
- 100 tuples per page,
- 1000 pages.

p_R

M

❖ Sailors:

- Each tuple is 50 bytes long,
- 80 tuples per page,
- 500 pages.

p_S

N

Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- ❖ In algebra: $R \bowtie S$. Common relational operation!
 - $R \bowtie S$ is large; $R \bowtie S$ followed by a selection is inefficient.
 - Must be carefully optimized.
- ❖ We will consider more complex join conditions later.
- ❖ *Cost metric*: # of I/Os. We will ignore output costs.

Simple Nested Loops Join

```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- ❖ For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
- Cost: $M + p_R * M * N = 1000 + 100 * 1000 * 500 = 1,000 + (5 * 10^7)$ I/Os.
- Assuming each I/O takes 10 ms, the join will take about 140 hours!

“Tuple at a time” Nested Loops Join

Page-Oriented Nested Loops Join

- ❖ For each *page* of R, get each *page* of S, and write out matching pairs of tuples $\langle r, s \rangle$, where r is in R-page and S is in S-page.
- **Cost:** $M + M * N = 1000 + 1000*500 = 501,000$ I/Os.
- Assuming each I/O takes 10 ms, the join will take about 1.4 hours.
- ❖ Choice of the *smaller* relation as the *outer*
- If smaller relation (S) is outer, cost = $500 + 500*1000 = 500,500$ I/Os.