

Midterm Review

CMPSCI 445
Spring 2018

Exam Thursday

- Please arrive promptly (or early if you can).
- Sit closer to door if you think you might finish early.
- Don't sit with your friends.
- No notes, books, electronic devices.
- No blue books.
- Just bring a pen or pencil.

A note on sample exam

- The SQL questions on the exam spend more time on basic SQL features, but ignore advanced stuff (data cubes, window functions, json)
- That advanced stuff is fair game. Expect some questions on it (like Homework 4)

Topics on the exam

Topic	Textbook Reference
Relational algebra	Ch 4
SQL review, SQL advanced	Ch 5 + slides from class
XML, JSON, XPath queries	27.6 + slides
Indexes and access methods	Ch 8, 9, 10, 11
Sorting and join algorithms	Ch 12, 13, 14

Review topics

- Relational algebra
- Nested queries
- Advanced SQL
- XML/XPath, JSON
- Access methods
- Sorting and Join Algorithms

Relational Algebra

Relational Algebra

- Operates on relations, i.e. *sets*
 - Later: we discuss how to extend this to *bags*
- Five basic operators:
 - Union: \cup
 - Difference: $-$
 - Selection: σ
 - Projection: Π
 - Cartesian Product: \times
- Derived or auxiliary operators:
 - Intersection, complement
 - Joins (natural, equi-join, theta join)
 - Renaming: ρ

Query equivalence

Definition: Query Equivalence

Two queries Q and Q' are equivalent if:

for all databases D , $Q(D) = Q'(D)$

Query Optimization

Is Based on Algebraic Equivalences

- Relational algebra has laws of commutativity, associativity, etc. that imply certain expressions are **equivalent**.
- They may be different in cost of evaluation!

$$\sigma_{c \wedge d}(R) \equiv \sigma_c(\sigma_d(R)) \quad \text{cascading selection}$$

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \quad \text{join associativity}$$

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S \quad \text{pushing selections}$$

- Query optimization finds the most efficient representation to evaluate (or one that's not bad)

Nested queries

Nested queries

- A **nested query** is a query with another query embedded within it.
- The embedded query is called the **subquery**.
- The subquery usually appears in the WHERE clause:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                   FROM Reserves R
                   WHERE R.bid = 103 )
```

(Subqueries also possible in FROM or HAVING clause.)

Correlated subquery

- If the inner subquery depends on tables mentioned in the outer query then it is a **correlated subquery**.
- In terms of conceptual evaluation, we must recompute subquery for each row of outer query.

```
SELECT  S.sname  
FROM    Sailors S  
WHERE   EXISTS ( SELECT *  
                  FROM    Reserves R  
                  WHERE   R.bid = 103  
                  AND R.sid = S.sid )
```

Correlation



Example

What does this query compute?

```
SELECT S.sid, S.name  
FROM   Sailors S  
WHERE  S.rating >= ALL (SELECT S2.rating  
                        FROM Sailors S2 )
```

- Find the sailors with the highest rating

exercise

- Emp(eid, ename, age, salary)
- Works(eid, did)
- Dept(did, dname, budget, managerid)

```
SELECT DISTINCT D.managerid
FROM Dept D
WHERE 1000000 < ALL (SELECT D2.budget
                     FROM Dept D2
                     WHERE D2.managerid = D.managerid )
```

Describe in English the result of this query

Example schema

- **Sales**(item-name, color, size, **number**)
 - **item-name**: {skirt, dress, shirt, pant}
 - **color**: {dark, pastel, white}
 - **size**: {small, medium, large}
 - **number**: number of units sold
- **Measure attributes**: measure some value; can be aggregated.
- **Dimension attributes**: define the dimensions on which measure attributes, and summaries of measure attributes, are viewed.

Cross Tabulation

or Cross-Tab

size: all		<i>color</i>			
<i>item-name</i>		dark	pastel	white	Total
	skirt	8	35	10	53
	dress	20	10	5	35
	shirt	14	7	28	49
	pant	20	2	5	27
	Total	62	54	48	164

- Values for one of the dimension attributes form the **row** headers
- Values for another dimension attribute form the **column** headers
- Other dimension attributes are listed on top (here, just *size*)
- Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.

Cross-tab as relation


size: **all**

<i>item-name</i>	<i>color</i>			
	dark	pastel	white	Total
skirt	8	35	10	53
dress	20	10	5	35
shirt	14	7	28	49
pant	20	2	5	27
Total	62	54	48	164

- SQL queries can only output relations.
- Cross-tabs can be represented as relations

<i>item-name</i>	<i>color</i>	<i>number</i>
skirt	dark	8
skirt	pastel	35
skirt	white	10
skirt	all	53
dress	dark	20
dress	pastel	10
dress	white	5
dress	all	35
shirt	dark	14
shirt	pastel	7
shirt	white	28
shirt	all	49
pant	dark	20
pant	pastel	2
pant	white	5
pant	all	27
all	dark	62
all	pastel	54
all	white	48
all	all	164

Data cube



- A data cube is a multidimensional generalization of a cross-tab
- In general, a data “cube” can have n dimensions; 3 shown above
- Cross-tabs can be used as views on a data cube

Window functions

- A **window function** performs a calculation across a set of table rows that are somehow related to the current row.
- This is related to the type of calculation that can be done with an aggregate function, but...
- Unlike regular aggregate functions, use of a window function does not cause rows to become grouped into a single output row. The rows retain their separate identities.
- Behind the scenes, the window function is able to access more than just the current row of the query result.

Window function syntax

- A window function call always contains an OVER clause directly following the window function's name and argument(s).
 - This is what syntactically distinguishes it from a regular function or aggregate function.
 - The OVER clause determines exactly how the rows of the query are split up for processing by the window function.
 - The PARTITION BY list within OVER specifies dividing the rows into groups, or partitions, that share the same values of the PARTITION BY expression(s).
 - For each row, the window function is computed across the rows that fall into the same partition as the current row.

Ranking

```
SELECT depname,  
       empno,  
       salary,  
       rank() OVER (PARTITION BY depname ORDER BY salary DESC)  
FROM empsalary;
```

depname	empno	salary	rank
develop	8	6000	1
develop	10	5200	2
develop	11	5200	2
develop	9	4500	4
develop	7	4200	5
personnel	2	3900	1
personnel	5	3500	2
sales	1	5000	1
sales	4	4800	2
sales	3	4800	2

Rank within current
row's partition

Order By is important

XML, XPath, and JSON

Relations converted to XML

STUDENT

sid	name	gender
1	Jill	F
2	Bo	M
3	Maya	F

Takes

sid	cid
1	445
1	483
3	435

COURSE

cid	title	sem
445	DB	F08
483	AI	S08
435	Arch	F08

One possible
representation:

```
<students>
  <student>
    <name>Jill</name>
    <gender>F</gender>
    <courses>
      <course cid=445>
        <title>DB</title>
        <sem>F08</sem>
      </course>
      <course cid=483>
        <title>AI</title>
        <sem>S08</sem>
      </course>
    </courses>
  </student>
  <student>
    ...
  </student>
</students>
```


Xpath: Summary

bib	matches a bib element
*	matches any element
/	matches the root element
/bib	matches a bib element under root
bib/paper	matches a paper in bib
bib//paper	matches a paper in bib , at any depth
//paper	matches a paper at any depth
paper book	matches a paper or a book
@price	matches a price attribute
bib/book/@price	matches price attribute in book , in bib
bib/book[./@price<"55"]/ author/lastname	matches...

I-clicker exercise

What does the following XPath expression return on the XML document below

`//A[A]/*/*/text()`




- A. [empty set]
- B. one two three four
- C. two three
- D. one four seven
- E. two three five six

answer on next slide

I-clicker exercise

What does the following XPath expression return on the XML document below

`//A[A]/*/*/text()`



- A. [empty set]
- B. one two three four
- C. two three
- D. one four seven
- E. two three five six

Comparing File Organizations

- ❖ **Heap files** (random order; insert at eof)
- ❖ **Sorted files**, sorted on $\langle age, sal \rangle$
- ❖ **Clustered B+ tree file**, Alternative (1), search key $\langle age, sal \rangle$
- ❖ Heap file with **unclustered B + tree index** on search key $\langle age, sal \rangle$
- ❖ Heap file with **unclustered hash index** on search key $\langle age, sal \rangle$

Cost Model for Our Analysis

We ignore CPU costs, for simplicity:

- **B:** The number of data pages
- **R:** Number of records per page
- **D:** (Average) time to read or write disk page
- Measuring number of page I/O's ignores gains of pre-fetching a sequence of pages; thus, even I/O cost is only approximated.
- Average-case analysis; based on several simplistic assumptions.

☞ *Good enough to show the overall trends!*

Operations to Compare

- ❖ Scan: Fetch all records from disk
- ❖ Equality search
- ❖ Range selection
- ❖ Insert a record
- ❖ Delete a record

Assumptions in Our Analysis

❖ Heap Files:

- Equality selection on key; exactly one match.

❖ Sorted Files:

- Files compacted after deletions.

❖ Indexes:

- Alt (2), (3): data entry size = 10% size of record
- Hash: No overflow chains.
 - 80% page occupancy => File size = 1.25 data size
- B+Tree:
 - 67% occupancy (typical): implies file size = 1.5 data size
 - Balanced with fanout F (133 typical) at each non-level

Assumptions (contd.)

❖ Scans:

- Leaf levels of a tree-index are chained.
- Index data-entries plus actual file scanned for unclustered indexes.

❖ Range searches:

- We use tree indexes to restrict the set of data records fetched, but ignore hash indexes.

Cost of Operations

	Scan	Equality	Range	Insert	Delete
Heap File	BD	.5BD	BD	2D	Search + D
Sorted File	BD	$D \log_2 B$	$D(\log_2 B + \text{\#matching pages})$	Search + BD	Search + BD
Clustered Tree Index	1.5BD	$D \log_F 1.5B$	$D(\log_F 1.5B + \text{\#matching pages})$	Search + D	Search + D
Unclustered Tree Index	$BD(R + .15)$	$D(1 + \log_F .15B)$	$D(\log_F .15B + \text{\#matching recs})$	Search + 3D	Search + 3D
Unclustered Hash Index	$BD(R + .125)$	2D	BD	4D	4D


➡ *Several assumptions underlie these (rough) estimates!*

General External Merge Sort

➡ *More than 3 buffer pages. How can we utilize them?*

❖ To sort a file with N pages using B buffer pages:

- Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.
- Pass 2, ..., etc.: merge $B-1$ runs.




Cost of External Merge Sort

$$1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$$

- ❖ Number of passes:
- ❖ Cost = $2N * (\text{\# of passes})$
- ❖ E.g., with 5 buffer pages, to sort 108 page file:
 - Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - Pass 2: 2 sorted runs, 80 pages and 28 pages
 - Pass 3: Sorted file of 108 pages

Block Nested Loops Join

- ❖ Take the smaller relation, say R, as outer, the other as inner.
- ❖ Use one buffer for scanning the inner S, one buffer for output, and use all remaining buffers to hold ``block'' of outer R.
- For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result.
- Then read next page in S, until S is finished.



Examples of Block Nested Loops

- ❖ Cost: Scan of outer + #outer blocks * scan of inner
 - #outer blocks = $\lceil \# \text{ pages of outer} / \text{block size} \rceil$
 - Given available buffer size B, block size is at most B-2.
 - $M + N * \lceil M / B-2 \rceil$
- ❖ With Sailors (S) as outer, let block be 100 pages of S:
 - Cost of scanning S is 500 I/Os; a total of 5 *blocks*.
 - Per block of S, we scan Reserves; 5*1000 I/Os.
 - Total = 500 + 5 * 1000 = 5,500 I/Os.
 - (a little over 1 minute)

Index Nested Loops Join

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

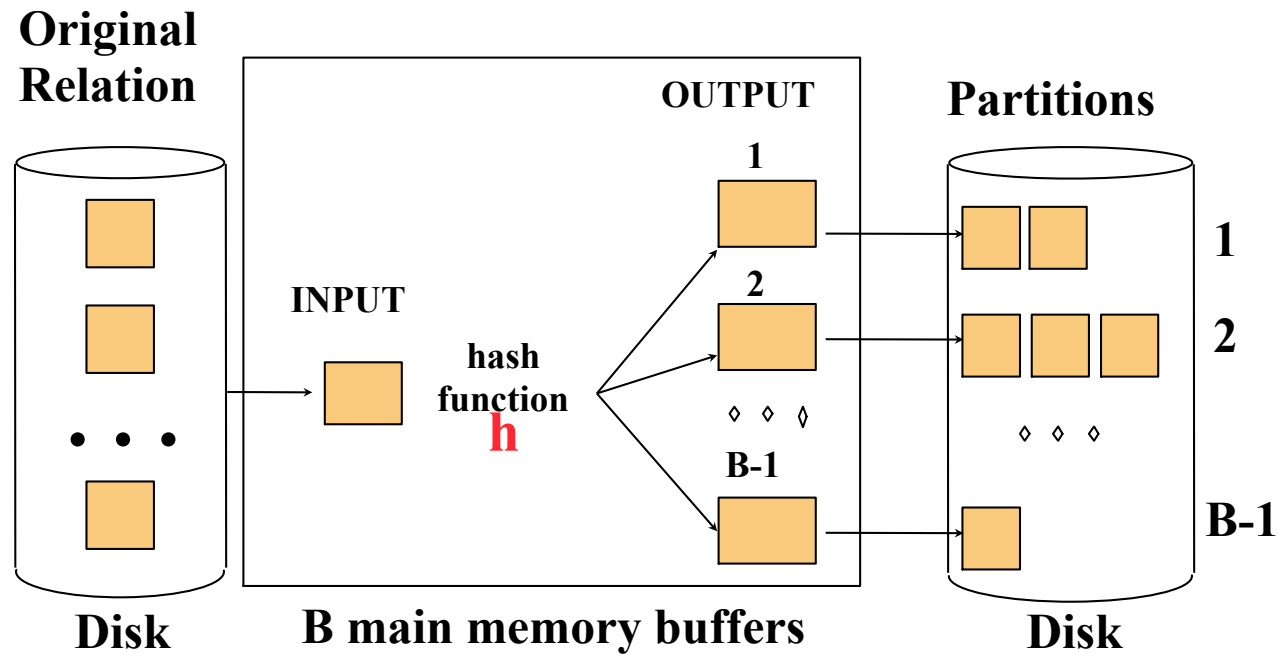
- ❖ If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
 - Cost: $M + (M * p_R) * \text{cost of finding matching S tuples}$
- ❖ For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples depends on clustering.
 - Clustered index: 1 I/O (typical).
 - Unclustered: up to 1 I/O per matching S tuple.

Sort-Merge Join ($R \bowtie_{i=j} S$)

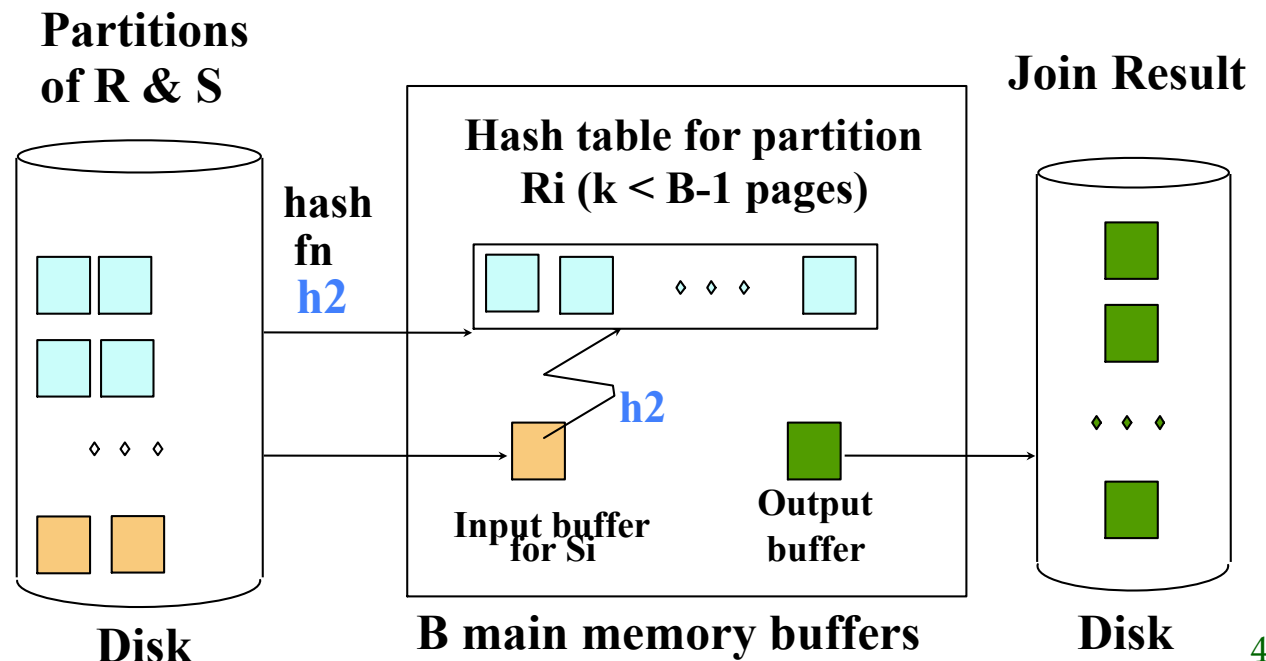
- ❖ (1) Sort R and S on the join column, (2) Merge them (on join col.), and output result tuples.
- ❖ Merge: repeat until either R or S is finished
 - *Scanning*: Advance scan of R until current R-tuple \geq current S tuple, advance scan of S until current S-tuple \geq current R tuple; do this until **current R tuple = current S tuple**.
 - *Matching*: Now all R tuples with same value in R_i (*current R group*) and all S tuples with same value in S_j (*current S group*) match; output $\langle r, s \rangle$ for all pairs of such tuples.
- ❖ R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)

Hash-Join

❖ Partitioning: Partition both relations using hash fn **h**: R tuples in partition *i* will only match S tuples in partition *i*.



❖ Probing: Read in partition *i* of R, build hash table on R_i using **h2** (\neq **h**!). Scan partition *i* of S, search for matches.



Observations on Hash-Join

- ❖ # partitions $\leq B-1$, and size of largest partition $\leq B-2$ to be held in memory. Assuming uniformly sized partitions, we get:
 - $M / (B-1) < (B-2)$, i.e., B must be $> \sqrt{M}$
 - Hash-join works if the smaller relation satisfies above.
- ❖ If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.
- ❖ If hash function h does not partition uniformly, one or more R partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this R -partition with corresponding S -partition.

Cost of Hash-Join

- ❖ Partitioning reads+writes both relns; $2(M+N)$. Probing reads both relns; $M+N$ I/Os. The total is $3(M+N)$.
 - In our running example, a total of 4500 I/Os using hash join, less than 1 min (compared to 140 hours w. NLJ).
- ❖ Sort-Merge Join vs. Hash Join:
 - With optimizations to Sort-Merge (not discussed in class), the cost is similar $\sim 3(M+N)$
 - Hash Join superior if relation sizes differ greatly.
 - Hash Join has been shown to be highly parallelizable.
 - Sort-Merge less sensitive to data skew; result is sorted.