

Introduction to Transaction Management

CMPSCI 445

Concurrency Control

- ❖ Concurrent execution of user programs is essential for good DBMS performance
- ❖ We must also cope with partial operations
- ❖ The **transaction** is the foundation for:
 - Concurrent execution
 - Recovery from system failure, incomplete ops

What is a Transaction?

- ❖ A **transaction** is the DBMS's abstract view of a user program: a sequence of reads and writes.

A simple transaction

- ❖ Imagine a simple banking application
 - Two database objects:
 - **A: balance of account A**
 - **B: balance of account B**
- ❖ Transaction T1:
 - “Transfer \$100 from account B to account A”.

T1: Transfer
Begin
$A = A + 100$
$B = B - 100$
End

The ACID Properties

- ❖ Database systems ensure the **ACID** properties:
 - Atomicity
 - Consistency
 - Isolation
 - Durability

Atomicity

- ❖ A very important property guaranteed by the DBMS for all transactions is that they are **atomic**.
 - User can think of a Xact as executing **all its actions** in one step, or executing **no actions at all**.
 - DBMS logs all actions so that it can undo the actions of aborted transactions.
- ❖ If it succeeds, the effects of write operations persist (**commit**);
- ❖ If it fails, no effects of write operations persist (**abort**)

Consistency

- ❖ Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins.
 - DBMS will enforce some ICs, depending on the ICs declared in CREATE TABLE statements.
 - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
- ❖ In banking example, sum ($A + B$) should be unchanged by execution.

Isolation

- ❖ Many concurrent transactions are running at one time.
- ❖ Each transaction should be **isolated** from the effects of other transactions
- ❖ Transactions should not be exposed to intermediate states created by other transactions.
- ❖ The net effect of concurrently running {T1 and T2 and T3} is equivalent to some serial order
 - No guarantee which serial order

Durability

- ❖ If transaction completes, its effects will persist in the database.
- ❖ In particular, if the system crashes before effects are written to disk, they will be **redone**
- ❖ Recovery manager is responsible for this.

The ACID Properties

- ❖ Database systems ensure the **ACID** properties:
 - **Atomicity**: all operations of transaction reflected properly in database, or none are.
 - **Consistency**: each transaction in isolation keeps the database in a consistent state (this is the responsibility of the user).
 - **Isolation**: should be able to understand what's going on by considering each separate transaction independently.
 - **Durability**: updates stay in the DBMS!!!

Two transactions

- “Transfer \$100 from account B to account A”
- “Add 6% interest to accounts A and B”

T1: Transfer
Begin
$A = A + 100$
$B = B - 100$
End

T2: Interest
Begin
$A = 1.06 * A$
$B = 1.06 * B$
End

Serial execution: T1, then T2

- Starting balances
 - $A = 1000$
 - $B = 2000$
- Execute T1
 - $A = 1100$
 - $B = 1900$
- Execute T2
 - $A = 1166$
 - $B = 2014$

T1: Transfer

Begin

$A = A + 100$

$B = B - 100$

End

T2: Interest

Begin

$A = 1.06 * A$

$B = 1.06 * B$

End

Serial execution: T2, then T1

- Starting balances
 - $A = 1000$
 - $B = 2000$
- Execute T2
 - $A = 1060$
 - $B = 2120$
- Execute T1
 - $A = 1160$
 - $B = 2020$

T2: Interest

Begin

$A = 1.06 * A$

$B = 1.06 * B$

End

T1: Transfer

Begin

$A = A + 100$

$B = B - 100$

End

Interleaved execution

❖ What other results are possible if operations of T1 and T2 are interleaved?

- Starting balances

- $A = 1000$
- $B = 2000$

T1: Transfer
...
$A = A + 100$
...
$B = B - 100$
...

T2: Interest
...
$A = 1.06 * A$
...
$B = 1.06 * B$
...

Interleaving operations

T1: Transfer	T2: Interest
$A = A + 100$	$A = 1.06 * A$
$B = B - 100$	$B = 1.06 * B$

Is this interleaving okay?

Interleaving operations

T1: Transfer	T2: Interest
$A = A + 100$ $B = B - 100$	$A = 1.06 * A$ $B = 1.06 * B$

How about this interleaving?

Goal: interleaved execution, with serial effects

- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. However, the net effect *must* be equivalent to these two transactions running **serially** in some order.

Scheduling Transactions

- ❖ A transaction is seen by DBMS as sequence of reads and writes
 - read of object O denoted $R(O)$
 - write of object O denoted $W(O)$
 - must end with Abort or Commit
- ❖ A **schedule** of a set of transactions is a list of all actions where order of two actions from any transaction must match order in that transaction.

A schedule

T1: Transfer	T2: Interest
$A = A + 100$	$A = 1.06 * A$ $B = 1.06 * B$
$B = B - 100$	



T1: Transfer	T2: Interest
Read(A) Write(A)	
	Read(A) Write(A) Read(B) Write(B)
Read(B) Write(B)	

Scheduling Transactions

- ❖ *Serial schedule*: Schedule that does not interleave the actions of different transactions.
- ❖ *Equivalent schedules*: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- ❖ *Serializable schedule*: A schedule that is equivalent to some serial execution of the transactions.
- ❖

Serializable Schedule

T1	T2
R(A) W(A)	
	R(A) W(A)
R(B) W(B)	
	R(B) W(B)
	Commit
Commit	

When can actions be re-ordered?

- ❖ Let I, J be two consecutive actions of T1 and T2
 - I=Read(O), J=Read(O)
 - I=Read(O), J=Write(O)
 - I=Write(O), J=Read(O)
 - I=Write(O), J=Write(O)
- ❖ If I and J are both reads, then they can be freely reordered.
- ❖ In all other cases, order impacts outcome of schedule.

Conflicting operations

- ❖ Two operations **conflict** if:
 - they operate on the same data object, and
 - at least one is a WRITE.
- ❖ Schedule outcome is determined by order of the conflicting operations.

Conflict Serializable Schedules

- ❖ Two schedules are **conflict equivalent** if:
 - Involve the same actions of the same transactions
 - Every pair of conflicting actions (of committed trans) are ordered the same way.
 - Alternatively: S can be transformed to S' by swaps of non-conflicting actions.
- ❖ Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule

Every conflict serializable schedule is serializable.

(exception: dynamic databases)

Conflict-serializable schedule

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)
	Commit
Commit	

Not conflict-serializable

T1	T2
R(A) W(A)	R(A) W(A) R(B) W(B) Commit
R(B) W(B) Commit	

Precedence graphs

- ❖ Directed graph derived from schedule S:
 - Vertex for each transaction
 - Edge from T_i to T_j if:
 - T_i executes Write(O) before T_j executes Read(O)
 - T_i executes Read(O) before T_j executes Write(O)
 - T_i executes Write(O) before T_j executes Write(O)

If edge $T_i \rightarrow T_j$ appears in precedence graph, then in any serial schedule equivalent to S, T_i must appear before T_j .

Dependency Graph

- ❖ Theorem: A schedule is **conflict serializable** if and only if its dependency graph is acyclic.

(A serializable order can be found by topological sort of the dependency graph.)


Construct precedence graphs:

T1	T2
R(A) W(A)	
	R(A) W(A)
R(B) W(B)	
	R(B) W(B) Commit
Commit	



Conflict serializable

T1	T2
R(A) W(A)	
	R(A) W(A) R(B) W(B) Commit
R(B) W(B) Commit	



Non-conflict serializable

Construct precedence graph:

T1	T2	T3
R(A)	W(A) Commit	
W(A) Commit		W(A) Commit