

# *Crash Recovery*

CMPSCI 445

Gerome Miklau

# *Review: the ACID Properties*

- ❖ Database systems ensure the **ACID** properties:
  - **Atomicity**: all operations of transaction reflected properly in database, or none are.
  - **Consistency**: each transaction in isolation keeps the database in a consistent state (this is the responsibility of the user).
  - **Isolation**: should be able to understand what's going on by considering each separate transaction independently.
  - **Durability**: updates stay in the DBMS!!!

# *Types of failure*

## ❖ Transaction failure

- partially-executed transaction cannot commit
- ➔ changes must be removed: ROLLBACK


## ❖ System failure

- volatile memory lost
- ➔ updates of committed Xact persist
- ➔ updates of aborted or partial Xacts removed

## ❖ Media failure

- corrupted storage media
- ➔ database brought up-to-date using backup

# Motivation



- ❖ Desired Behavior after system restarts:
  - T1, T2 & T3 should be **durable**.
  - T4 & T5 should be **aborted** (effects not seen).

# *Undo and Redo*

## ❖ UNDO:

- removing effects of incomplete or aborted transaction (for atomicity)

## ❖ REDO:

- re-instating effects of committed transactions (for durability)

- ❖ The work the recovery subsystem must do to support UNDO and REDO depends on **key policies** of the buffer manager.

# *More on Steal and Force*

- ❖ STEAL (why enforcing Atomicity is hard)
  - *To steal frame F*: Current page in F (say P) is written to disk; some Xact holds lock on P.
    - What if the Xact with the lock on P aborts?
    - Must remember the old value of P at steal time (to support **UNDO**ing the write to page P).
- ❖ NO FORCE (why enforcing Durability is hard)
  - What if system crashes before a modified page is written to disk?
  - Write as little as possible, in a convenient place, at commit time, to support **REDO**ing modifications.

# *Handling the Buffer Pool*

- ❖ **Force** every write to disk?
  - Poor response time.
  - But provides durability.
- ❖ **Steal** buffer-pool frames from uncommitted Xacts?
  - If not, poor throughput.
  - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Recovery <b>Trivial</b>	
No Force		<b>Desired</b>

# Basic Idea: Logging



- ❖ Record REDO and UNDO information, for every update, in a *log*.
  - Sequential writes to log (put it on a separate disk).
  - Minimal info (diff) written to log, so multiple updates fit in a single log page.
- ❖ Log: An ordered list of REDO/UNDO actions
  - Log record contains:
    - Before image (for UNDO), After image (for REDO)
  - and additional control info (which we'll see soon).

<XID, pageID, offset, length, old data, new data>



# Write-Ahead Logging (WAL)

- ❖ The Write-Ahead Logging Protocol:

- ① Must force the **log record** for an update before the corresponding **data page** is overwritten on disk.
- ② Must **write all log records** for a Xact before commit.

- ❖ #1 guarantees Atomicity.

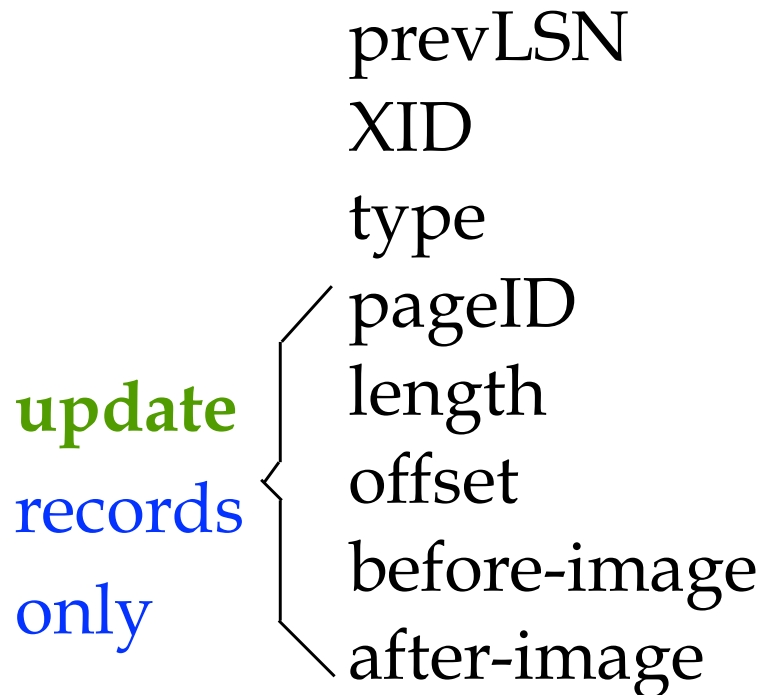
- ❖ #2 guarantees Durability.

- ❖ Exactly how is logging and recovery done?

- the ARIES algorithm (we won't see details of this in this class)

# *Log Records*

## LogRecord fields:



Possible log record types:

- ❖ **Update**
- ❖ **Commit**
- ❖ **Abort**
- ❖ **End** (signifies end of commit or abort)
- ❖ **Compensation Log Records (CLRs)**
  - for UNDO actions

# *The Big Picture: What's Stored Where*




## **LogRecords**

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



**Data pages**  
each  
with a  
pageLSN  
**master record**



## **Xact Table**


lastLSN  
status

## **Dirty Page Table**

recLSN

**flushedLSN**

# Crash Recovery: Big Picture



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - **Analysis**: Figure out which Xacts committed since checkpoint, which failed.
  - **REDO** *all* actions.
    - ◆ (repeat history)
  - **UNDO** effects of failed Xacts.

# *Crash during recovery*

- ❖ Crashes during UNDO handled by logging CLRs
- ❖ What happens if system crashes during Analysis or Redo?
  - Analysis: all work is lost, but analysis begins again.
  - Redo: Just redo again -- redo idempotent. Some pages may have been written to disk before crash but this will be evident.

# *Summary of Logging/Recovery*

- ❖ **Recovery Manager** guarantees Atomicity & Durability.
- ❖ Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.